Thyleads Operations Playbook

v2.0 — April 2026 Owner: CEO | Maintained by: Business Head | Reviewed quarterly Audience: All Thyleads roles — CEO, Business Head, Pod Leads, GTM Engineers, Account Managers, Data Analysts

---

How this document is organized

This is the master operations playbook. It covers the complete client lifecycle from contract signature through Deal Momentum, the role-by-role responsibilities that make it work, the tool stack and how to use it effectively, and the dashboard architecture that enforces the standards. It is written to be read in sequence on first read and used as a reference document thereafter.

Structure:

1. Operating principles
2. Roles & responsibilities (CEO, Business Head, Pod Lead, GTME, AM, DA)
3. Tool stack & how to use it effectively
4. Pre-onboarding (Day -7 to 0)
5. Track A — Infrastructure (Days 1–21, parallel)
6. Track B — Onboarding (Days 1–14, main track)
7. Phase 5 — Active campaign operations (Day 15+)
8. Phase 6 — Deal Momentum (post-first-meeting)
9. Quality gates summary
10. Dashboard workflow spec (for dev)
11. Failure modes & escalation
12. Implementation rollout across pods
13. Appendices

---

1. Operating principles

These are the non-negotiables that hold across every pod, every client, every campaign.

The dashboard is the source of truth. No data lives in Google Sheets after Day 0. Every account, contact, score, signal, email variant, send log, reply, meeting, transcript, and Deal Momentum touch is captured in the Thyleads dashboard. If it isn't logged, it didn't happen.

Quality gates are non-negotiable. Each phase ends with a gate. A pod cannot progress to the next phase until the gate is signed off. Skipped gates create downstream failures that cost three to five times more to fix.

Humans own judgment. AI owns volume. AI agents draft the "why now" hook, generate content variants, classify replies, and surface pattern recommendations. Humans decide which accounts to target, approve content before it sends, qualify positive replies, and run client conversations.

Compliance is a Day 0 default. RFC 8058 one-click unsubscribe, plain visible unsubscribe in footer, CAN-SPAM compliance, and DMARC at p=quarantine minimum are baked into infrastructure setup before the first email sends.

Pilot before scale, every time. Every campaign launches as a 100–200 contact pilot. Mailbox reputation, content performance, and persona fit are validated against pilot data before production scale.

Deal Momentum is not optional. Every meeting that books triggers a structured 60-day follow-up cadence. The first meeting is the start of the deal cycle, not the end of Thyleads' work.

Variance is the enemy of retention. Two clients in two different pods should experience the same quality of work. Pod-level autonomy is welcome on tactics; pod-level autonomy on standards is not.

---

2. Roles & responsibilities

Each role is described with mandate (what they own), daily/weekly cadence (how they spend time), reporting line (who they report to), and key metrics (how performance is measured).

CEO (Akash)

Mandate. Vision, positioning, brand, hiring/firing decisions, strategic partnerships, fundraising, and final authority on playbook revisions. The CEO is not in the day-to-day client work but reviews the operating system that runs it.

Cadence. Weekly: review aggregate dashboard metrics across all pods (target attainment, retention risk, gap closure rate). Monthly: Repository pattern review with Business Head — what's the system learning, what should change. Quarterly: playbook revision based on accumulated learnings, hiring plan against pipeline.

Reporting line. Reports to investors and the board. Business Head reports to CEO.

Metrics. Aggregate retention rate, NPS across active clients, gross margin per pod, playbook adherence index, pipeline-to-revenue conversion across the agency.

When to involve the CEO. Client churn risk that has escalated past Business Head; potential lighthouse client wins; strategic disputes between pods on resource allocation; quarterly playbook revisions; any compliance event that materially threatens the business.

Business Head

Mandate. Operational excellence across all four pods. Owns the Gaps & Risks register, SOP enforcement, cross-pod resource allocation, and is the escalation point for retention risk before it

reaches the CEO. The Business Head is the chief operating officer of Thyleads in everything but title.

Cadence. Daily: review pod health dashboard, scan for SLA breaches and gate slips. Weekly: 1:1 with each Pod Lead, review of all 12 clients' health status, gap triage meeting (Wednesdays), all-hands ops standup (Mondays). Monthly: pod performance reviews, repository pattern review with CEO, hiring forecast review.

Reporting line. Reports to CEO. Pod Leads report to Business Head.

Metrics. Gap closure velocity, SOP adherence index across pods, pod performance variance (lower is better), retention rate, cross-pod knowledge transfer (Repository contributions per pod per month).

Owned artifacts. Gaps & Risks register (Outbound tab), Repository (institutional memory), pod-level performance dashboards, SOP version control, all-hands ops meeting agenda.

Pod Lead

Mandate. Quality and outcomes for the 3–4 clients in their pod. Owns Gates 3, 4, and 5 — the pod-internal quality gates that prevent broken work from reaching clients. Coaches GTME, AM, and DA on craft. First escalation point for any client issue.

Cadence. Daily: pod standup (15 minutes), review reply triage queue, review SLA breaches. Per-client weekly: 30-minute review with the AM ahead of the client review call. Per-client monthly: deeper QBR-style review with all pod members, with explicit retention assessment. Weekly with Business Head.

Reporting line. Reports to Business Head. GTME, AM, and DA in their pod report to the Pod Lead for line management; functional escalation goes through them.

Metrics. Pod target attainment, client retention within pod, content QA pass rate, gate adherence rate, time-to-first-meeting for new clients, NPS within pod's clients.

Authority. Approves Gate 3 (content + infra sign-off), Gate 4 (production scale), Gate 5 (Deal Momentum monthly review). Can pause campaigns. Can reassign work within pod. Cannot hire/fire (escalates to Business Head).

GTM Engineer

Mandate. Everything strategic and technical between client kickoff and live campaign: discovery, ICP, account universe, scoring, persona research, content, infrastructure setup and monitoring, campaign execution, pilot data analysis. The GTME is the strategist + builder of every client engagement.

Cadence. New client (Days 1–14): primary lead, full focus. Active client (Day 15+): 60–70% of time on optimization, content iteration, infra monitoring; remainder on new client onboarding rotations. Weekly: pod standup, content review with Pod Lead, repository contribution. Daily: mailbox health check (5 min in Infrastructure tab), campaign performance scan (10 min in Campaigns tab).

Reporting line. Reports to Pod Lead.

Metrics. Reply rate per campaign, positive reply rate, mailbox health score (no mailbox at risk for >24hr), content QA pass rate (first-submission), pilot-to-production conversion rate, time-to-first-meeting.

Owned artifacts. Kickoff brief, ICP brief, account universe + scoring, persona briefs, all content variants, infrastructure registry, pilot analysis reports.

Account Manager

Mandate. Client relationship from Day 1, reply triage and qualification, meeting coordination, Deal Momentum cadence ownership, weekly client reporting. The AM is the human face of Thyleads to the client and to the prospect once a positive reply lands.

Cadence. Continuous reply monitoring during business hours (replies SLA is 4 hours). Daily: morning Deal Momentum cadence review (which prospects need touches today), client AE check-ins via Slack/WhatsApp where applicable. Weekly: client review call (30 minutes per client), pod standup, internal AM huddle for cross-learning. Per meeting: pre-meeting prep doc, post-meeting outcome capture, Deal Momentum cadence kickoff.

Reporting line. Reports to Pod Lead.

Metrics. Reply SLA adherence (4hr), positive reply qualification rate, meeting show rate, 1st-meeting → 2nd-meeting conversion, Deal Momentum cadence adherence rate, client retention contribution.

Owned artifacts. Reply triage queue, qualification notes, meeting prep docs, post-meeting outcome cards, Deal Momentum activity log, weekly client reports.

Data Analyst

Mandate. Stakeholder mapping for accounts that automated tooling cannot handle. The DA is an exception-handler, not a manual labor pool — the role evolves once Apollo, Clay, Crustdata, and Sales Nav are pulling automated stakeholders. The DA fills the 20–30% residual gap with manual research and verifies data quality.

Cadence. Per-project basis. When a GTME hands off accounts, the DA has 5 working days to map. Daily: pull queue from dashboard, work through assigned accounts. Weekly: QA sample of own work (5%), pod standup.

Reporting line. Reports to Pod Lead for line management; takes project-level direction from GTME.

Metrics. Account coverage rate (target: 95%+ within 5 working days), stakeholder accuracy rate (target: 95%+ via verification agent and 5% manual QA), exception completion rate (accounts automation couldn't handle).

What the DA does NOT do. Does not contact prospects. Does not write or send emails. Does not handle replies. Does not enrich emails or phone numbers (that's GTME's Clay workflow). Does not communicate with clients.

3. Tool stack & how to use it effectively

The stack is opinionated. Pods do not pick their own tools; the tools listed below are the standard. Adding a new tool to the stack requires Business Head approval and a 30-day pilot.

Outbound execution

Smartlead ($200/mo currently). Email campaign execution and warming. Used for all email sequences across all clients.

How to use it well. One Smartlead campaign per persona × use-case combination, never one mega-campaign. A/B testing is set up at the subject-line level on every campaign. Mailbox rotation is on (sends spread across all healthy mailboxes for that client). Daily send caps stay under 30 per mailbox in the first 30 days, scaling to 50 only after sustained green health metrics. Throttle-based sending (15-second random delay between emails) is the default, not bulk burst.

Common pitfall. Teams often run too few campaigns with too many variants stuffed inside, making it impossible to read which variant is actually working. Always one persona × use-case = one campaign.

Integration with our stack. API connected to Thyleads dashboard; reply data flows into the Replies tab; mailbox health metrics flow into the Infrastructure tab; campaign performance flows into the Campaigns tab.

Heyreach ($100/mo for 3 seats — 4th pod gap to fix). LinkedIn campaign execution.

How to use it well. Connection requests stay under 80 per day per LinkedIn account to avoid restrictions. Warm-up before any outreach: 2 weeks of engagement-only behavior (likes, comments) on target persona content. Sequencing: connection request → wait 2 days → message #1 (no pitch, just curiosity) → wait 4 days → message #2 (light value) → wait 5 days → message #3 (soft CTA). LinkedIn voice notes outperform text on follow-up touches for senior personas.

Common pitfall. Using the same LinkedIn account across multiple clients. Each client needs a dedicated LinkedIn persona where possible. Account restrictions hit hardest when LinkedIn detects multi-client patterns.

Integration. API into Thyleads dashboard; LinkedIn replies flow into the same Replies tab as Smartlead.

Data & enrichment

Apollo ($220/mo currently for 1–2 seats; recommended upgrade to 4 seats — one per pod). Primary contact and company database. Used for initial account universe pulls and bulk stakeholder mapping.

How to use it well. Build saved search filters per ICP pattern (e.g., "MarTech VP Marketing 200-2000 employees India"). Use the "Buying Intent" filter for warm-tier targeting. Email verification

within Apollo is okay for triage but not for sending — always re-verify with ZeroBounce before pushing to Smartlead.

Common pitfall. Treating Apollo data as fresh. Apollo data ages quickly; titles change, people leave, domains get acquired. Always run a freshness check (verification agent) before campaigns launch.

Crustdata. Alternative B2B database with better international coverage, especially for India-specific firmographics, hiring signals, and tech stack data.

How to use it well. Use Crustdata for signal-based account discovery — "companies that hired 5+ growth marketers in last 90 days," "companies that switched CRM in last 6 months," "companies with $5M+ Series B in last quarter." This is where the "why now" signals come from. Crustdata is stronger than Apollo for India market depth.

Common pitfall. Trying to use Crustdata as the primary contact database — it's better at company-level signals than person-level depth. Pair it with Apollo or Sales Nav for stakeholder mapping.

Clay (free credits via mailbox stacking — recommend $149/mo paid as insurance). Orchestration and enrichment layer that ties multiple data sources together.

How to use it well. Clay is best as the connective tissue — pulling raw lists from Apollo or Crustdata, enriching emails through waterfall (Apollo → Hunter → LeadMagic), verifying with ZeroBounce, scoring with custom signals, and pushing clean data into Smartlead. Build reusable Clay tables per ICP pattern. Use AI columns for the "why now" snippet generation (this is where Opus 4.7 plugs in).

Common pitfall. Free credit limits hit fast at scale. The mailbox-stacking trick works until Clay enforces it. Budget for paid before that happens.

Integration. Clay tables read from and write to Thyleads dashboard via webhook.

LinkedIn Sales Navigator. Manual research, stakeholder verification, persona language sourcing.

How to use it well. Sales Nav is best for verification ("does this person actually still work here?") and for the deep persona research where the GTME reads 20 LinkedIn posts from a persona to absorb their language. Build saved searches per persona. Use TeamLink to identify warm intros that the client AE could leverage.

Common pitfall. Treating Sales Nav as a primary outreach tool. It's a research tool. Outreach goes through Heyreach.

Hunter.io / LeadMagic. Email finding tools used as fallback in the Clay enrichment waterfall.

How to use them well. These run after Apollo doesn't return an email. Always run found emails through ZeroBounce before sending. Hunter is stronger for US/EU; LeadMagic has better India coverage.

ZeroBounce. Email verification — last step before sends.

How to use it well. Verify every email before it enters Smartlead. Anything that comes back as "catch-all" or "unknown" gets dropped from the campaign, not sent to. Bounce rates above 3% will get mailboxes deplatformed; ZeroBounce is the firewall against that.

Infrastructure

Google Workspace / Outlook 365. Mailbox provisioning. Roughly 60/40 split between Workspace and Outlook is recommended for diversification — if Gmail tightens its policy, Outlook mailboxes stay alive (and vice versa).

How to use them well. Provision mailboxes with realistic-looking names and full profiles (photo, signature, LinkedIn link). Empty/anonymous mailboxes get filtered as spam by Gmail's ML. Each mailbox costs $6–12/month — at 12 clients × 12 mailboxes that's $864–1700/month in pure mailbox cost. Budget accordingly.

GoDaddy / Namecheap. Domain registration. Default: 3–5 alt-domains per client.

How to use them well. Domain naming matters. "ThyXClient.com" gets flagged faster than realistic alternatives like "client-team.co" or "client-digital.io". Avoid keywords like "marketing," "sales," "outbound" in domains — these correlate with cold outreach and get classified faster. Domain age helps; consider keeping a "domain bank" of pre-aged domains for new clients.

Google Postmaster Tools + Microsoft SNDS. Free deliverability monitoring direct from Gmail and Outlook.

How to use them well. Connect every domain on Day 1 of warming. Check daily for the first 30 days, weekly thereafter. Postmaster shows spam rate (must stay under 0.1%), domain reputation (must stay High), and IP reputation. Any drop triggers an automated mailbox pause via the Infrastructure tab.

MXToolbox. Free DNS and blacklist checking.

How to use it well. Run weekly against every domain. If a domain shows up on a blacklist (Spamhaus, SORBS, etc.), pause sends from that domain immediately and run delisting requests. Most blacklists delist within 48 hours if abuse stops.

Meeting intelligence

Google Calendar. Source of truth for all meetings. Both Thyleads and client AE on every external meeting.

Fireflies.ai. Meeting recording, transcription, AI summary. Already in use for CleverTap; rolling out across all clients.

How to use it well. Fireflies bot auto-joins every meeting where Thyleads or the client AE is invited. Configure the AI summary to extract the structured fields: attendees actual vs invited, outcome category, objections raised, competitors mentioned, use cases that resonated, next steps, sentiment. Webhook the output into the Thyleads dashboard. Set custom fields specifically for our outcome taxonomy.

Common pitfall. Some enterprise prospects (BFSI, Healthcare) refuse bots. Have a fallback: AM takes structured notes during the call into a templated form, runs Whisper-based transcription on a recording where allowed, or uses Otter.ai for audio-only. Document the exception in the meeting record.

AI layer

Claude Opus 4.7. Orchestration, content generation, judgment-heavy tasks, transcript analysis, pattern synthesis.

How to use it well. One agent = one job. Don't build mega-prompts that try to do five things; build composable agents. Use system prompts to encode the Thyleads style and constraints (e.g., "always 50–90 words for first-touch email, never use the words 'just,' 'quick,' or 'circling back'"). Maintain a prompt registry in the dashboard so prompts version-control alongside the dashboard code.

Claude Haiku. High-volume classification — reply triage, content QA scoring, signal scanning. Roughly 1/15th the cost of Opus per token.

How to use it well. Anywhere you're scoring/classifying with a defined output schema, use Haiku. Anywhere you're generating prose or making judgment calls, use Opus. Reply triage in particular is a Haiku job — classifying inbound as positive/negative/OOO/wrong-person/auto-reply runs thousands of times a day and doesn't need Opus.

Approximate monthly API spend at 12-client steady state: $300–600.

Internal

Thyleads dashboard (Next.js). System of record. Everything ladders into this.

Slack. Team communications, real-time SLA breach alerts, pod standups via huddle.

Notion. Living version of this playbook (until the dashboard absorbs SOP enforcement). Repository documentation. Onboarding documents for new hires.

WhatsApp. Used for Deal Momentum communication with prospects who have engaged (never for cold outreach), and for client AE coordination where the client uses WhatsApp as primary channel — common for India market.

---

4. Pre-onboarding (Day -7 to Day 0)

The clock does not start at kickoff. The moment a contract is signed, Track A (infrastructure) begins. The client doesn't need to know — but the warming clock matters.

Day -7 (contract signed). Business Head assigns client to a pod based on capacity, vertical match, and DA bandwidth. Pod Lead is notified. GTME is identified. AM is assigned and immediately begins context-gathering on the client's space (reading their website, watching their last conference talk, reading their last 10 LinkedIn posts).

Day -7 to -5. Domain procurement: GTME procures 3–5 alt-domains via GoDaddy or Namecheap. Domain naming follows the pattern guidance above. DNS setup begins (A records, MX records).

Day -4 to -1. Mailbox provisioning: 3 mailboxes per domain via Google Workspace or Outlook 365. Authentication setup: SPF, DKIM, DMARC at p=quarantine. Branded tracking subdomain configured. Postmaster Tools and SNDS connected. Compliance items: List-Unsubscribe-Post header configured at Smartlead level, plain unsubscribe link template in mailbox signatures.

Day 0. Kickoff call scheduled with client. By this point, mailboxes are connected to Smartlead warming and have started accumulating warming days. By the time the pilot launches on Day 11 of Track B, mailboxes have ~18 days of warming on them — within the safe zone.

---

5. Track A — Infrastructure (Days 1–21, parallel to onboarding)

This track runs alongside the discovery/targeting/build/calibrate phases. The GTME owns infrastructure but spends only ~30 minutes a day on it during this period (daily health check + responses to anomalies).

Day	Activity	Owner
1	Mailboxes connected to Smartlead warming	GTME
1	Postmaster + SNDS connected for every domain	GTME
1–21	Daily health check via Infrastructure tab (5 min)	GTME
7	Mid-warming review: any mailbox below threshold pauses + investigation	GTME + Pod Lead
14	Pre-pilot infra readiness check (warmed mailboxes >14 days, all auth green, no blacklist)	GTME
21	Full warming complete; all mailboxes at production capacity	GTME

Mailbox health thresholds (auto-pause triggers). Spam complaint rate > 0.1%, bounce rate > 3%, Postmaster reputation drops below "High," domain shows on any major blacklist (Spamhaus, SORBS), Smartlead warmup score below 80%. Any single trigger pauses that mailbox automatically and notifies the GTME via Slack.

Compliance checklist (must be green before any pilot send).

- SPF record published and validating
- DKIM signing active and validating
- DMARC at p=quarantine minimum, monitoring at p=reject within 30 days

- List-Unsubscribe-Post header (RFC 8058) on every send
- Plain visible unsubscribe link in footer of every variant
- Tracking subdomain branded (not Smartlead default)
- Postmaster Tools and SNDS connected
- DNC list and existing-customer suppression run on contact list

---

6. Track B — Onboarding (Days 1–14, main track)

Phase 1 — Discovery (Days 1–3)

Day 1 — Kickoff & context capture. 60-minute call with the client lead (Founder, CRO, or Head of Sales). GTME leads, AM joins. Capture product overview, current GTM state, three best customers and why they bought, three lost deals and why, competitive positioning, brand do's and don'ts. Output: Kickoff Brief in dashboard. Client receives written summary within 4 hours.

Day 2 — Customer & AE interviews. GTME conducts 30-minute calls with 2–3 of the client's actual customers + 1 of the client's AEs. This is the highest-leverage research of the entire onboarding. Direct quotes go into the language bank for content drafting.

Day 3 — ICP & value-prop synthesis. GTME synthesizes into a one-page ICP brief: refined ICP with disqualifiers, 2–3 personas with pain narratives, competitive positioning, 3–5 messaging angles, top 3 "buy now" signals. **Gate 1 — Client signs ICP brief.** 48hr SLA; escalation to Pod Lead at 48hr, Business Head at 72hr.

Phase 2 — Targeting (Days 4–7)

Day 4 — Account universe. GTME pulls from Apollo + Clay + Crustdata + Sales Nav. Volume = 2.5–3x monthly meeting target in account count. Suppression flags applied. Logged in dashboard, not Sheets.

Day 5 — Account scoring. Every account scored 0–100 against the rubric (base fit 30, signals 40, engagement 20, why-now 10). AI Research Agent generates "why now" snippet per top-tier account. **Gate 2 — Top 200–500 scored accounts approved by client.** 48hr SLA.

Day 6 — DA handoff. Approved accounts handed to DA with brief: account list, required job titles, stakeholder count target (3–5 per account), 5 working day deadline. Automated stakeholder pull runs first; DA fills the residual gap. Verification agent validates titles. Target: 95%+ coverage.

Day 7 — Persona research deep-dive. GTME builds persona briefs: day-in-the-life, top 5 pains in their words, top 3 KPIs, common objections, communication patterns. AI Research Agent surfaces relevant content from podcasts, LinkedIn, Reddit.

Phase 3 — Build (Days 8–11)

Day 8 — Skeleton selection. Pull 2–3 skeletons from Repository (hiring-signal, pain-point, use-case, competitor-switch). Per skeleton: 3 persona angles × 3 use-case angles = 9–12 first-touch

variants per campaign. Cadence: 5–7 touches over 14–21 days, mixing email and LinkedIn.

Day 9 — Content drafting. AI Content Agent drafts variants from persona briefs and language bank. GTME edits, never blindly accepts. Subject-line A/B variants set up.

Day 10 — QA & sign-off. AI QA Agent scans every variant for spam words, mobile readability, personalization placeholder validation, framework compliance (Pain → Solution → Social Proof → CTA), length (50–90 words). Pod Lead reviews and signs off in dashboard. **Gate 3 — Pod Lead signs off content + infra readiness.** Launch button disabled in UI until both checkboxes signed.

Day 11 — Pilot launch. 100–200 contacts, evenly distributed across persona × use-case cells. 1–2 mailboxes only for the pilot. Sends spread over 3 days. AM begins active reply monitoring; 4-hour SLA on positive replies activates.

Phase 4 — Calibrate (Days 12–14)

Day 12 — Pilot data analysis. Read by variant, never aggregate. Metrics: delivered, opened, replied, positive reply, mailbox health. AI Targeting Feedback Agent surfaces patterns.

Day 13 — Iterate. Kill bottom-performing variants. v2 round on top performers. Adjust scoring rubric if needed. Update Repository.

Day 14 — Production launch. Gate 4 — Pod Lead signs off on production scale. If pilot data is strong, scale to full mailbox set + full account list. If weak, loop back to Day 9. Weekly client review locked (default: Friday 30 minutes). AM takes over as primary client interface.

---

7. Phase 5 — Active campaign operations (Day 15+)

Once production launches, the engagement enters steady state. Each week looks like:

Daily. GTME does morning mailbox health check (5 min) and campaign performance scan (10 min). AM does reply triage continuously through business hours. AM runs Deal Momentum touches per cadence (see Phase 6).

Weekly. Pod standup Monday morning. Internal pod review of all clients on Wednesday. Client review call (30 min) end of week. AM delivers weekly client report 2 hours before the client review call.

Monthly. GTME contributes to Repository: what worked, what didn't, with attribution. Pod Lead runs deeper QBR-style review with all pod members. Business Head reviews pod-level metrics and gap closure rate. Optimization round on content based on accumulated data.

Weekly client report contents. Sends, deliveries, opens, replies, positive replies (as raw and as rate). Meetings booked, meetings completed, meetings to 2nd meeting. Active Deal Momentum cadences and stage. Mailbox health summary. Top-performing variants. Recommended next-week experiments. Repository contributions made.

---

8. Phase 6 — Deal Momentum (post-first-meeting)

Deal Momentum is what separates Thyleads from agencies that drop off after the first meeting. Every meeting that books triggers a structured 60-day cadence owned by the AM, with content support from the GTME, in coordination with the client AE.

The cadence assumes the prospect is qualified and the meeting was productive. If the meeting was unqualified, AM marks the lead "disqualified" in the dashboard and the cadence does not run.

Pre-meeting (T-2 to T-0)

T-2 days. AM sends meeting prep doc to client AE: prospect company, attendees, why this meeting was booked (what email/LinkedIn message they responded to), recent signals from the account (hiring, funding, news), suggested talking points, suggested objection-handling, prospect's role and likely KPIs. AM ensures Fireflies bot is invited.

T-0 (meeting day). AM confirms with client AE 2 hours before meeting. Fireflies joins.

Meeting + immediate post

T-0 (meeting completes). AI Meeting Intelligence Agent processes Fireflies transcript within 5 minutes. Outputs structured fields: attendees actual vs invited, outcome category (qualified / disqualified / needs-nurture / pilot-interest / closed-now), objections raised, competitors mentioned, use cases that resonated, stakeholders referenced, agreed next steps, sentiment.

T-0 + 30 min. AM reviews AI summary, edits if needed, dispatches structured outcome card to client AE via email or Slack. Card format: what happened in 3 bullets, what the prospect said in their words, recommended next 3 touches, any flags for the AE.

Cadence

T+1 day. Thank-you + recap email. Drafted by AI from transcript, edited by AM, sent from client AE's address (or co-signed). Includes: meeting recap in 3 lines, agreed next steps, one piece of relevant content (case study, ROI doc, etc.).

T+3 days. Value content drop. Specific to the objection raised or use case discussed. AI suggests best-fit content from the client's library; AM selects and personalizes. For India market, this is also where a WhatsApp touch becomes appropriate ("hi, sharing the case study you asked about during our chat").

T+7 days. Stakeholder expansion ask. "Who else from your team should see this? Happy to set up a brief overview for them." Goal: identify and engage other decision-makers. AM updates stakeholder map in dashboard.

T+14 days. Vertical-specific case study. Tailored to prospect's industry (FinTech case study to a FinTech prospect, HRTech to HRTech). Reinforces social proof.

T+21 days. Explicit next-step ask. "Ready to discuss the proposal?" or "Should we set up a tech deep-dive with your engineering lead?" Forces a yes/no/not-yet response.

T+30 days. Close-the-loop. "Should we keep this opportunity active or pause for now?" Re-qualifies intent. If response is "pause," the prospect moves to nurture (signal-triggered re-engagement).

T+45 to T+60 days. Reactivation triggers. AI Targeting Feedback Agent monitors for new signals on the prospect's account: new funding, hiring, leadership change, content engagement, competitor switches. Any trigger reactivates the cadence with a fresh "why now."

Channel mix per touch

Touch	Primary channel	Backup channel
T+0 outcome card	Email/Slack to AE	—
T+1 thank-you	Email	LinkedIn
T+3 value content	Email	WhatsApp (India) / LinkedIn
T+7 stakeholder ask	Email	LinkedIn
T+14 case study	Email	LinkedIn
T+21 next-step	Email	Phone (AM) / WhatsApp
T+30 close-the-loop	Phone (AM)	Email
T+45–60 reactivation	Email + LinkedIn	WhatsApp

Quality gate for Deal Momentum

Gate 5 — Monthly Deal Momentum review. Pod Lead reviews every Deal Momentum cadence at month-end. Metrics: cadence adherence rate (touches sent on schedule), 1st-meeting → 2nd-meeting conversion, time-from-meeting-to-next-step, win rate by cadence touch (which touch typically converts). Cadences with adherence below 90% trigger AM coaching. Cadences with conversion below pod average trigger content review with GTME.

What the client AE owns vs what AM owns

The client AE is the deal owner; the AM is the deal accelerator. The AE makes pricing decisions, runs technical deep-dives, and signs the contract. The AM keeps the deal warm in between AE touches, coordinates schedules, ensures follow-ups happen on cadence, and surfaces objections back to the AE early.

AE-owned: pricing conversations, contract negotiation, technical scoping, demo customization, executive sponsorship calls.

AM-owned: cadence adherence, meeting coordination, content delivery, stakeholder map maintenance, signal monitoring, objection-tracking, weekly status to AE.

9. Quality gates summary

Gate	When	Owner	What it enforces
Gate 1	Day 3	Client	ICP brief signed before targeting begins
Gate 2	Day 5	Client	Top accounts approved before contact mining
Gate 3	Day 10	Pod Lead	Content + infra ready before pilot send
Gate 4	Day 14	Pod Lead	Pilot data validated before production scale
Gate 5	Monthly	Pod Lead	Deal Momentum cadence adherence + conversion

---

10. Dashboard workflow spec (for dev)

The dashboard must enforce the playbook. SOPs that aren't enforced by the system get skipped under deadline pressure.

Client onboarding state machine. Field `client_onboarding_state`: values `pre_onboarding`, `discovery`, `targeting`, `build`, `calibrate`, `production`. State transitions are gated:

- `pre_onboarding → discovery` — kickoff scheduled (auto on calendar event)
- `discovery → targeting` — Gate 1 signed (client signature timestamp captured)
- `targeting → build` — Gate 2 signed (client account approval timestamp)
- `build → calibrate` — Gate 3 signed (Pod Lead checkbox on content + infra)
- `calibrate → production` — Gate 4 signed (Pod Lead production approval)

Each gate creates an audit log entry with timestamp, signer, notes.

Parallel infrastructure state. Independent of onboarding state. Field `infra_state`: `provisioning`, `warming`, `ready`, `at_risk`, `paused`. Per-mailbox health metrics roll up to per-domain, then per-client.

Required fields per state. "What's blocking next state" panel surfaces missing fields automatically.

SLAs hardcoded.

- 48 hours from ICP brief send to client signature (escalation at 48hr/72hr)
- 4 hours from positive reply received to AM acknowledgment
- 5 working days from DA handoff to stakeholder mapping completion
- 30 minutes from meeting completion to outcome card dispatch

- 24 hours per Deal Momentum touch (cadence adherence)

Auto-progression rules. Where signals are unambiguous, state advances automatically. Where judgment is required (pilot review, outcome categorization), manual sign-off.

Account universe schema. Company name, domain, ICP fit flag, score (0–100), score breakdown by dimension, "why now" snippet, source, suppression flags, client approval status, last touch date, current campaign assignment.

Contact schema. First name, last name, LinkedIn URL, verified title, persona tag, email (with enrichment status + verification timestamp), phone, DA owner, verification status, last touch date.

Meeting auto-logging. Calendar event with both Thyleads + client AE → Fireflies bot auto-join → placeholder record in dashboard linked to original campaign + reply → post-meeting webhook → AI summary populates structured fields → AM gets 30-min SLA notification.

Deal Momentum module. Per-meeting cadence record with auto-scheduled touches at T+1, T+3, T+7, T+14, T+21, T+30, T+45/60. AI drafts each touch. AM reviews/edits/sends. Status visible to client AE in client-facing dashboard.

Repository auto-publish. Every completed campaign auto-publishes a record: client (anonymized cross-client), persona × use case, performance metrics, top variants, what worked. Searchable by ICP/persona/industry for new client onboarding.

---

11. Failure modes & escalation

Gate 1 misses by 48 hours. Pod Lead intervenes with client. If no signature within 72 hours, Business Head joins. Most often this is sponsor context drift; a 30-min alignment call resolves it.

Pilot data is weak (Gate 4 blocked). Loop back to Day 9 with refined skeletons. Do not scale "to see what happens." Cost of broken production is mailbox reputation damage that takes weeks to recover.

DA stakeholder mapping below 80% by Day 11. Reassign to a stronger DA in the pod or Pod Lead runs partial automation backfill. Do not launch with thin coverage — pilot data will be unreadable.

Mailbox health degrades during pilot or production. Infrastructure tab auto-pauses affected mailbox and routes to warming. AM and GTME notified within 15 minutes. Pilot/production continues on remaining healthy mailboxes only. GTME investigates root cause within 24 hours.

Client AE unresponsive to Deal Momentum. AM raises in weekly review with explicit ask. If no improvement after 2 weeks, Pod Lead joins client review. Business Head escalation at month 2 if pattern persists. This is a retention risk signal, not a tactical issue.

Compliance event (blacklist, spam complaint spike, GDPR/CAN-SPAM violation). GTME pauses campaign immediately, notifies Pod Lead and Business Head within 1 hour. Business

Head decides on disclosure to client (default: yes, transparently). Root cause fix before any restart.

Client churn signal. Defined as: 3+ consecutive missed targets, NPS drop below 7, client AE disengagement on Deal Momentum, or explicit dissatisfaction. Pod Lead owns retention plan within 48 hours; Business Head reviews; CEO involved if churn is likely.

---

12. Implementation rollout across pods

This playbook is rolled out in three waves to manage change risk:

Wave 1 (Weeks 1–2). Pilot with one pod, ideally Pod 4 (Thyleads internal — current performance baseline 65% completion is the easiest to A/B against). One new client onboarded entirely under this playbook. Existing clients in that pod continue under current process (no disruption).

Wave 2 (Weeks 3–6). All new client onboardings across all four pods follow this playbook. Existing clients begin transitioning to Deal Momentum cadence and Repository contribution. Weekly Business Head reviews against playbook adherence.

Wave 3 (Weeks 7–12). Existing clients fully migrated to playbook. Dashboard build complete (state machine, gates, schemas, auto-logging, Deal Momentum module). Repository populated with first 90 days of pattern data. Quarterly playbook revision begins.

Success metrics for the rollout. Time-to-first-meeting drops by 30%, content QA pass rate (first submission) above 80%, mailbox health no incidents above 24hr, DA completion rate above 95%, Deal Momentum cadence adherence above 90%, retention rate improves quarter-over-quarter.

---

13. Appendices

Appendix A — Kickoff call agenda

60 minutes, structured, no drift. Minute 0–10: product walkthrough by client. Minute 10–20: current GTM, what's working, what's broken, why outsource. Minute 20–35: best customers — three real names, why they bought, what they almost bought instead, what would have stopped them. Minute 35–45: lost deals — three real names, what killed each, what the winning competitor said. Minute 45–55: competitive positioning — top 3 competitors, where you win, where you lose, lies they tell. Minute 55–60: brand do's and don'ts.

Appendix B — Customer interview script

Six questions in order, do not skip ahead. Why did you buy? What were you using before, and what made you switch? What pain pushed you to act this quarter and not the previous one? Who else needed to say yes — what did you say to them? What almost stopped you? If you had to describe what this product does in one sentence to someone in your role, what would you say? The verbatim language from question 6 is gold.

Appendix C — Pod Lead Gate 3 sign-off checklist

All variants pass AI QA Agent scoring. Personalization placeholders validated on 10 random rows. List-Unsubscribe-Post header confirmed in test send. Plain visible unsubscribe in footer. SPF/DKIM/DMARC green on all sending mailboxes. Mailbox warming minimum 14 days. Tracking subdomain branded. Content matches ICP brief and persona pain narratives. No competitor mentioned negatively. Brand do's and don'ts respected.

Appendix D — Deal Momentum cadence template per touch

T+1 thank-you. Subject: "Notes from our chat — [topic]." Body: 3-line recap of meeting, agreed next steps, one piece of relevant content. Signed by client AE.

T+3 value drop. Subject: "[Specific objection or use case] — quick read." Body: framing of what was discussed, content link with 1-line context, soft "happy to walk through this if useful." Signed by AM (or AE if relationship-driven).

T+7 stakeholder ask. Subject: "Helpful for the broader team?" Body: acknowledgment of what was discussed, ask for who else should see this, offer to set up a brief overview. Signed by AM.

T+14 case study. Subject: "How [similar company] solved [similar problem]." Body: 1-paragraph context, case study link, soft connection to their situation. Signed by AM.

T+21 next-step. Subject: "Where should we take this from here?" Body: recap progress, explicit yes/no/not-yet ask on proposal or next meeting. Signed by AE.

T+30 close-the-loop. Phone call from AM to prospect or AE. If no answer, voicemail + email follow-up. Honest re-qualification.

T+45–60 reactivation. Subject driven by signal: "Saw you raised your Series B — congrats. Wanted to reconnect on [topic]." Body: signal acknowledgment, fresh "why now," soft ask.

Appendix E — RACI matrix summary

Activity	CEO	Business Head	Pod Lead	GTME	AM	DA
Vision & positioning	A	C	I	I	I	I
Playbook revision	A	R	C	C	C	I
Gap closure	I	A	R	C	C	I
Client onboarding	I	I	A	R	C	C
ICP & targeting	I	I	A	R	I	I
Stakeholder mapping	I	I	A	C	I	R
Content drafting	I	I	A	R	C	I
Infrastructure	I	I	A	R	I	I

Activity	CEO	Business Head	Pod Lead	GTME	AM	DA
Reply triage & qualification	I	I	A	C	R	I
Meeting coordination	I	I	A	I	R	I
Deal Momentum	I	C	A	C	R	I
Weekly client reporting	I	C	A	C	R	I
Retention	I	A	R	C	C	I

R = Responsible (does the work) | A = Accountable (owns the outcome) | C = Consulted | I = Informed

---

End of v2.0. This playbook is a living document. The Repository captures what each new client teaches us, and the playbook evolves quarterly based on patterns surfaced by the Targeting Feedback Agent and field feedback from pods.